

Vol 0 Chapter 4 — The Isotropy Group: Where the Symmetries of Physics Come From

Keeper (author pass)

2026-05-23 Saturday

Chapter 4 — The Isotropy Group: Where the Symmetries of Physics Come From

If you take standard physics' organizational structure as given, one of the things that has to be explained somewhere is *why it organizes the way it does*. There are spatial rotations and Lorentz boosts. There are gauge symmetries — electric charge, color, weak isospin. There are discrete symmetries — parity, time reversal, charge conjugation. There are conservation laws, one for each continuous symmetry, and a famous theorem (CPT) tying the discrete symmetries together. None of these are derived in standard physics; each is taken as a structural input, observed in nature and built into the theory by hand.

In BST, all of them are consequences of one fact about D_{IV}^5 : the structure of the group that *stabilizes a point of the substrate*. Pick any point of D_{IV}^5 — the origin of the bounded-domain realization is the conventional choice — and ask which elements of the substrate's symmetry group $SO_0(5, 2)$ leave that point fixed. The answer is a smaller group, called the isotropy subgroup at the point. For D_{IV}^5 , the isotropy subgroup factors:

$$\text{Isotropy} = SO(5) \times SO(2) \text{ plus a Möbius involution.}$$

Each of the three pieces of this factorization is responsible for a specific class of physical symmetry. The $SO(5)$ piece organizes spacetime observables — position, momentum, angular momentum, spin. The $SO(2)$ piece organizes internal symmetries — electric charge and chirality. The Möbius involution is responsible for parity, and the way it interacts with the weak sector explains why parity is *violated* exactly where it is violated. The full symmetry group $SO_0(5, 2)$, larger than the isotropy, supplies the Hamiltonian, Lorentz invariance, and the CPT theorem.

This chapter unpacks all of that. The result will be a single picture: the standard list of physical symmetries is what the isotropy of D_{IV}^5 produces, in the order in which it produces them.

4.1 The dimensions add up

Before assigning physical meaning, the bookkeeping is worth doing once.

The substrate's full symmetry group $SO_0(5, 2)$ has dimension 21 (this is the dimension of the orthogonal Lie algebra in signature $(5, 2)$). The isotropy subgroup $SO(5) \times SO(2)$ has dimension $10 + 1 = 11$. The difference, $21 - 11 = 10$, is the real dimension of the symmetric space D_{IV}^5 — the number of “directions in which one can move away from a fixed point.”

This last number is important. The isotropy subgroup fixes a point; the *complement* of the isotropy in $SO_0(5, 2)$ — what mathematicians call the coset directions, written $\mathfrak{m}$ — moves the point. Position operators and momentum operators live in the coset directions, not in the isotropy. Angular momentum and spin live in the isotropy. This distinction is the structural source of the difference between “where you are” (a coset displacement) and “how you are oriented” (an isotropy rotation). We will use it explicitly in §4.2.

The Möbius involution is one additional discrete element — not a Lie-group factor but a single $\mathbb{Z}_2$ generator that the rank-2 structure of D_{IV}^5 forces. Lyra’s T2443 result (May 21, 2026) established this: rank = 2 on a Type IV domain implies a $\text{Pin}(2)$ double cover of the isotropy, and the $\text{Pin}(2)$ double cover carries an additional discrete element that we will identify as the parity operation.

So: three pieces. $SO(5)$, $SO(2)$, and a Möbius $\mathbb{Z}_2$. Each will produce its own slice of physics.

4.2 The $SO(5)$ piece: position, momentum, angular momentum, spin

The $SO(5)$ factor of the isotropy is the rotation group of five-dimensional Euclidean space. It has $5 \cdot 4/2 = 10$ independent generators, all of them rotations (skew-symmetric, no translations).

The first thing to be careful about — and it is the kind of thing that has caught the team out in the past, so it is worth slowing down here — is that position and momentum operators are not in $SO(5)$. They are in the *coset complement* $\mathfrak{m}$ to the isotropy. The reason is operational. A position operator is supposed to measure “where on the substrate you are,” and “where” is by definition a displacement from a reference point. Displacements move the base point; they are not isotropy transformations. So position lives in the coset, not in $SO(5)$.

Concretely, the Lie algebra of $SO_0(5, 2)$ decomposes as

$$\mathfrak{so}(5, 2) = (\mathfrak{so}(5) \oplus \mathfrak{so}(2)) \oplus \mathfrak{m},$$

where $\mathfrak{m}$ has real dimension 10. This 10-dimensional coset complement, exponentiated, generates the substrate’s displacement structure. Position $\hat{X}$ is realized on the Bergman Hilbert space $H^2(D_{IV}^5)$ as multiplication by the coordinate z ; momentum $\hat{P}$ is realized as the Wirtinger derivative $-i\partial_z$. These are operators built from the coset directions, and they satisfy the canonical commutation relation $[\hat{X}, \hat{P}] = i\hbar$ — which, in BST, is not an axiom but a *consequence* of the Bergman kernel’s reproducing property (the explicit derivation is Lyra T2419 and T2422).

The $SO(5)$ factor itself supplies the operators that *do* have rotational character:

- Angular momentum $\hat{L}$ comes from the 10 generators of $SO(5)$ acting on Bergman space functions. The familiar three-dimensional angular momentum of laboratory physics is the restriction to the $SO(3) \subset SO(5)$ subgroup that rotates the three physical spatial dimensions. (We will see in §4.5 why “three” — it traces back to $N_c = 3$.)
- Spin $\hat{S}$ comes from the K-type structure of the substrate’s representations under $SO(5) \times SO(2)$. Internal spin is not an additional structure tacked onto orbital angular momentum; it is the substrate’s intrinsic representation index, half-integral or integral depending on whether the $\text{Pin}(2)$ double cover acts non-trivially or trivially on the K-type — which in turn depends on the bosonic-versus-fermionic character of the state. This is the structural origin of the spin-statistics distinction, which we will treat properly in Volume 5; for now, what matters is that spin lives where it does because the substrate’s isotropy has a $\text{Pin}(2)$ double cover, and the double cover has fermionic and bosonic sectors.

The five operators that live in the $SO(5)$ sector — position, momentum, angular momentum, spin, plus parity by the route described below — are exactly the spacetime-side observables of standard quantum mechanics. None of them was postulated; each is a structural consequence of the isotropy decomposition.

4.3 The $SO(2)$ piece: electric charge and chirality

The $SO(2)$ factor of the isotropy is one-dimensional and acts as a phase rotation. Equivalently, it is a $U(1)$ symmetry. The substrate’s natural action is multiplication by $e^{i\alpha Q}$ on substrate states, with Q the eigenvalue of the $SO(2)$ generator — the substrate’s *internal* phase.

What this generator means physically depends on what it acts on.

On scalar substrate states, the $SO(2)$ generator is electric charge. Particles with $SO(2)$ weight ± 1 are charged; weight 0 is neutral. Antiparticles have opposite weight from particles. Hadrons, being substrate states with composite color structure inherited from the $N_c = 3$ substrate sub-structure (Volume 2’s subject), carry fractional

charges — the famous $\pm 1/3, \pm 2/3$ of quark electric charges. The quantization in units of $1/N_c$ is not a postulate; it falls out of the $SO(2)$ representation theory crossed with the color structure. We will derive this carefully in Volume 2.

On spinor substrate states, the $SO(2)$ generator acts as a chiral phase rather than a charge multiplication. Its eigenvalue is γ^5 , the chirality operator that distinguishes left-handed from right-handed particles. Same factor, different action — the difference between charge and chirality, in BST, reduces to whether the substrate state being acted on is in a scalar or spinor K-type of the substrate's representation theory.

The two operators have the same algebraic source. The hypercharge $U(1)_Y$ of the Standard Model is the substrate's $SO(2)$ generator; electromagnetism's $U(1)_{em}$ is the unbroken combination remaining after Weinberg mixing with $SU(2)$ weak isospin. Both descend from the substrate's $SO(2)$ isotropy factor, with the Weinberg mixing angle itself a substrate-derived quantity (Volume 2, Chapter 5).

4.4 The Möbius involution: parity, and why the weak sector violates it

The third piece of the isotropy is the Möbius involution. It is a single $\mathbb{Z}_2$ element — its square is the identity — that arises from the rank-2 $\text{Pin}(2)$ double cover. Geometrically, it is an orientation-reversing transformation of D_{IV}^5 that interchanges complex-conjugate structures. Operationally, lifted to the substrate's Hilbert space, it becomes the parity operator $\hat{P}$.

The parity operator has the properties we expect:

$$\hat{P}^2 = 1, \quad \hat{P}\hat{X}\hat{P} = -\hat{X}, \quad \hat{P}\hat{p}\hat{P} = -\hat{p}, \quad \hat{P}\hat{S}\hat{P} = +\hat{S}.$$

Position and momentum flip sign under parity, as in standard quantum mechanics; spin does not flip. Eigenstates of parity are even or odd; physical observables are decorated with definite parity.

The deeper question is why parity is *conserved* in some sectors and *violated* in others — strong and electromagnetic interactions preserve parity, the weak interaction does not. Standard physics treats this as a fact discovered experimentally (Wu's experiment, 1957) and inserted into the theory by hand. BST derives it.

The mechanism is what Casey's W-21 work calls Möbius locality. The Möbius involution is not a global symmetry of D_{IV}^5 in every sector. In sectors where the substrate's commitment dynamics commute with Möbius — the strong and electromagnetic sectors — parity is conserved. In the weak sector, the substrate's chirality structure (the $SO(2)$ action on spinors) is intrinsically chirally asymmetric: weak doublets pair left-handed components together, leaving right-handed components in singlets. This chirality asymmetry does *not* commute with Möbius. So in the weak sector, the substrate Hamiltonian and the Möbius involution have non-zero commutator, and parity is violated — exactly in this sector, exactly for this reason.

This is the kind of derivation BST is supposed to do. Standard physics observes that the weak sector violates parity. BST shows that the substrate's structure *requires* this violation, in the weak sector specifically, with the magnitude controlled by the chirality coupling.

4.5 The full $SO_0(5, 2)$ and what only it can provide

The isotropy gets us position, momentum, angular momentum, spin, charge, chirality, and parity — seven of standard physics' observables, with their commutation structure and (in parity's case) violation rule built in. The remaining observables require the *full* symmetry group $SO_0(5, 2)$, not just its isotropy subgroup.

The Hamiltonian. Time evolution is a one-parameter subgroup of $SO_0(5, 2)$ that is not contained in the isotropy. Elie's K52a Session 29 work established that the substrate Hamiltonian is the Casimir operator of $SO_0(5, 2)$ acting on the appropriate K-type — its lowest non-trivial eigenvalue is $C_2 = 6$, the BST primary integer we met in Chapter 2. The Hamiltonian therefore inherits a structural energy unit set by C_2 .

Lorentz invariance. The Lorentz group $SO_0(3, 1)$ of four-dimensional Minkowski spacetime is a subgroup of $SO_0(5, 2)$: it is the subgroup that preserves the four-dimensional conformal structure D_{IV}^5 inherits from its em-

bedding (Chapter 1, §1.3). So Lorentz invariance is not a separate axiom; it is what one inherits when restricting from the substrate’s full symmetry to its $(3 + 1)$ -dimensional spacetime sector.

The CPT theorem. Combining the Möbius involution (parity $\hat{P}$), the anti-unitary Klein operator (time reversal $\hat{T}$ — Lyra T2433), and the $SO(2)$ weight negation (charge conjugation $\hat{C}$ — Lyra T2434), the composite $\hat{C}\hat{P}\hat{T}$ is structurally an element of $SO_0(5, 2)$ ’s conformal action. Its action commutes with the substrate Hamiltonian *globally*, in every sector, for structural reasons. This is the CPT theorem of standard quantum field theory — Lüders 1954, Pauli 1955 — recovered from substrate structure rather than postulated.

The pattern is: the more refined the symmetry of standard physics, the more complete the group of $SO_0(5, 2)$ needed to derive it. Isotropy gets you the local observables. Full group gets you the time evolution, the relativistic structure, and the universal CPT result.

4.6 Why three spatial dimensions

A question that has bothered physicists for a long time is why physical space has three dimensions rather than two or four. BST’s answer is on display in this chapter, even though the question’s full treatment is in Volume 1.

The $SO(5)$ factor has five rotational dimensions. Under the substrate’s natural splitting they decompose as $5 = N_c + \text{rank} = 3 + 2$ — three spatial dimensions controlled by the BST primary $N_c = 3$ (this is where the three of “color charge” first appears), and two further dimensions controlled by rank. The three physical spatial dimensions visible in laboratory physics are the N_c piece. The remaining rank-2 piece supplies, via the $\text{Pin}(2)$ double cover, the spin-statistics structure of fermions and bosons — it is internal rather than spatial, even though it lives in the same $SO(5)$ factor of the isotropy.

So “why three spatial dimensions” reduces to “why $N_c = 3$,” which Chapter 2 derived from the Mersenne cascade ($N_c = M_{\text{rank}}$) and from the trefoil topology of confined color-bearing matter. The chain runs all the way back to the substrate’s structural integers.

4.7 What comes next

Chapter 5 examines boundary conditions: how the substrate’s interior dynamics, which we have now organized by the isotropy decomposition, connect to its Shilov boundary — the lower-dimensional structure where the four-phase cycle’s emission output is read off. The boundary is where many of BST’s sharpest experimental falsifiers live.

Chapter 6 returns to the integer web. We met it briefly in Chapter 2; Chapter 6 elevates it to a Casey-named structural principle and shows that the cross-identities among BST primary integers carry their own structural information, beyond what any individual integer encodes.

Chapter 7 — the operator zoo — collects together the operators we have introduced in this chapter (position, momentum, angular momentum, spin, charge, chirality, parity, Hamiltonian) plus a few more (time reversal, charge conjugation, Bell-CHSH, particle number), and shows the canonical decomposition under the isotropy factors. It is the operational reference chapter for everything in Volumes 1 and 2.

Where to look this up: The rank-2 $\text{Pin}(2)$ double cover and the parity origin of Möbius are Lyra T1925 and T2443 (May 2026). The time-reversal anti-unitary Klein operator is T2433; charge conjugation as $SO(2)$ weight negation is T2434. The position and momentum coset realizations are T2419 and T2422; angular momentum on $SO(5)$ is T2425. The chirality-spin-statistics result deriving from rank-2 $\text{Pin}(2)$ grading is Paper #133. The substrate Hamiltonian as $SO_0(5, 2)$ Casimir is Elie’s K52a Session 29. Möbius locality producing weak-sector parity violation is Casey’s W-21 work. Standard references on coset decompositions of Lie algebras and isotropy actions on symmetric spaces are again Helgason, *Differential Geometry, Lie Groups, and Symmetric Spaces*, Chapters IV–V. For the spin-statistics theorem in standard quantum field theory, Streater and Wightman’s *PCT, Spin and Statistics, and All That* (Princeton, 1964) remains the canonical text; BST’s substrate-level derivation is intended to be read alongside it, not as a replacement.