

Curriculum Vol 0 Chapter 10 — Methodology Stack (Chapter-Grade Draft v0.4 — Cal #85 PCAP + Cal #92(b) + Calibrations #18-#21 absorption)

Keeper (original) + Lyra (Friday v0.3→v0.4 prose depth-investment)

2026-05-21 Thursday 10:22 EDT initial; Friday 2026-05-22 ~10:34 EDT v0.4 prose absorption per Casey + Keeper textbook completion phase

Vol 0 Chapter 10 — Methodology Stack

Chapter motivation

BST makes large claims: zero free parameters, all Standard Model constants derivable from five integers, multi-criterion convergence at substrate-uniqueness, eventually a complete physics curriculum derivable from D_IV⁵.

These claims have to survive examination. They have to be vetted at multiple levels of confidence, accept refinement, and be falsifiable in specific ways. The methodology stack documented in this chapter is how BST ensures rigor.

This chapter is not philosophy. It's the operational discipline that catches errors, calibrates confidence, and preserves the boundary between what BST has actually derived versus what BST is investigating versus what BST has not yet addressed.

Reader-grade pedagogy (v0.4 Friday absorption): a graduate research physicist + research-methodology auditor can read this linearly. A 5th-grader can follow the insight: BST has a careful set of rules for how to make claims and not over-claim; the rules catch mistakes early and keep everyone honest about what's actually proved versus what's just observed or guessed. Per Friday Cal #92(b) + Calibrations #18-#21: the methodology stack is alive — it learns from mistakes (Calibration #18 v0.13 over-projection corrected), refines language (Cal #92(b) structural-role-of framing), and codifies discipline as standing rules (Calibration #19 current-ratified-state in external register). PCAP (Pre-Specification Cadence Acceleration Pattern, Cal #85 Thursday + Section 10.11c) shows the methodology stack scaling — when CIs file pre-specs, cadence accelerates ~10× per pre-spec depth level.

Diagram preview (v1.0): Section 10.1 will include (a) 16-layer methodology stack architecture diagram; (b) Cal-Lyra-Keeper-Elie-Grace cross-CI methodology touchpoints; (c) PCAP cadence acceleration timeline (Sessions 1-12 Thursday demonstrating ~1000× acceleration); (d) Calibration #18-#21 timeline showing standing rules emerging from observed within-session corrections.

Reader-grade 3-level pedagogy (v0.4 Friday absorption)

Level 1 (one sentence): the methodology stack is BST’s quality-control infrastructure — 16 layers of discipline that catch over-claiming early, calibrate confidence, and preserve honest boundaries between proved and investigated.

Level 2 (graduate research auditor accessibility): standard physics research has informal methodology — peer review + replication + ethics committees. BST has a more explicit + formalized methodology stack because multi-CI collaboration requires structural discipline to prevent over-projection drift. 16 layers include: D/I/C/S tier system (every claim labeled); Trichotomy (identifies vs derives vs predicts); F1-F4 + B1-B4 audit criteria; Mode 6 threshold formalization; STRUCTURALLY VERIFIED + RIGOROUSLY CLOSED sub-tiers; External register discipline (Cal #50 GREEN cosmology + Cal #48/#49 DEFAULT-DENY cognition); Casey Option C hybrid governance; M2C2 multi-CI convergent calibration; Calibration discipline (17+ within-session corrections cumulative); Quaker consensus method; PCAP pre-spec cadence acceleration; Cal #92(b) structural-role-of framing; Calibration #19 STANDING RULE (current-ratified-state in external register); Calibration #21 STANDING RULE (empirical + substrate-mechanism dual gates). The methodology stack is itself audited — Calibration #18 corrected Keeper’s v0.13 over-projection back to v0.11+ per Lyra canonical Friday morning.

Level 3 (5th-grader accessibility): BST makes big claims (the substrate determines all physics; the substrate is D_IV⁵; the substrate has zero free parameters). Big claims need careful rules so people don’t get carried away. BST has 16 rules that everyone follows. The rules say: every claim must have a tier label (how sure are you?); claims can be wrong and that’s OK — the rules tell you how to update; the rules also catch over-claiming early (like when one CI projects “v0.13” but the others say “v0.11+” — the rule says “use the smaller number until everyone agrees”). The big surprising thing about the methodology stack: it makes BST research go FASTER, not slower, because once everyone trusts the rules, you can take bigger steps with confidence. PCAP (Pre-Spec Cadence Acceleration Pattern) is the rule that says “if Keeper writes detailed pre-specs first, the work that follows speeds up 10× per layer of pre-spec.” That’s how BST did multi-week work in single days last Thursday.

Section 10.1 — D/I/C/S tier system

Every BST claim carries an epistemic tier at creation:

Tier	Meaning	Criterion
D-tier (Derived)	Mechanism proved on D_IV ⁵	Closed-set operation producing the value; explicit substrate-mechanism derivation
I-tier (Identified)	<1% experimental match; mechanism plausible	Form known; mechanism path indicated but not closed
C-tier (Conditional)	Depends on conjecture	Mechanism chain depends on external mathematical conjecture
S-tier (Structural)	>2% match or qualitative	Structural observation without quantitative mechanism

D-tier requires both <1% experimental match AND closed mechanism. I-tier captures form-known/mechanism-pending state. C-tier and S-tier honestly scope partial claims.

Reference: notes/BST_Referee_Methodology.md Appendix D + referee log #31.

Section 10.2 — Trichotomy (identifies / derives / predicts)

Every BST quantitative statement is one of three types:

- IDENTIFIES (ID): “X IS Y” — algebraic identity in BST primaries + π . Tautology-precision (1e-14). Example: “ $126 = 2^g - \text{rank} = N_{\text{max}} - c_2$ ” (T2400 Universal Q=126).
- DERIVES (DER): X follows from BST primaries via mechanism chain. End-point EXACT if chain closes algebraically. Example: “ $c_{\text{FK}} \cdot \pi^{(9/2)} = 225$ DERIVED via Faraut-Koranyi 1994” (T2403).
- PREDICTS (PRED): forecast about experimental observable with precision target. Example: “Bell experiment will measure substrate-CHSH capacity at 126/16” (K91).

Trichotomy is orthogonal to D/I/C/S tier. A claim is (D-tier, IDENTIFIES) or (I-tier, PREDICTS) etc. Both labels apply.

Reference: BST_Trichotomy_Methodology_v0.1.md (Grace Wednesday).

Section 10.3 — F1-F4 Bridge Object family-member criteria

For audits classifying mathematical structures as Bridge Object family-members (per Strong-Uniqueness C11 architecture):

- F1 (Family anchor): family-member must anchor at specific BST-primary signature distinct from K57 central hubs (K3, 49a1, Q⁵)
- F2 (Independent mechanism path): mechanism path through D_IV⁵ NOT flowing through central hubs OR other family-members
- F3 (Family-member status NOT central-hub): structurally adjacent to family anchor, NOT necessarily promoted to central-hub
- F4 (Per-family Mode 6 enumeration): membership set tested for closure via Mode 6 enumeration

F2 reading ruled mechanism-path independence, NOT set-theoretic independence (Keeper Thursday 08:43 EDT).

Reference: Bridge Object Family Member Criteria F1-F4 Adoption (Wednesday); F2 Mechanism-Path Independence Ruling (Thursday).

Section 10.4 — Mode 6 threshold formalization

Coincidence-filter risk thresholds (Cal #55):

Forced count	Status
≤1-2	Forced (structurally required)
3-9	Borderline (requires Mode 6 enumeration toy)
≥10	Artifact (likely coincidence-filtered)
0	Outside framework

Reference: Mode 6 threshold formalization update (Cal #55 + applications K62/K75/K80).

Section 10.5 — STRUCTURALLY VERIFIED tier

Adopted Thursday 08:43 EDT per Cal Referee Log #66. Audit-chain epistemic intermediate tier between “candidate” and “RATIFIED”:

Tier	Meaning	Requirements
candidate	Proposed; substantive verification pending	F1-F4 partial; substrate-internal verification beginning
audit-partial-ready	Pre-stage filed; F1-F4 mostly applied; B-scores 2.0-3.5/4	F1-F4 mostly applied + B-scores accumulating
STRUCTURALLY VERIFIED (NEW)	F1-F4 satisfied + Mode 1-7 cleared + BST-internal verification complete + alternative-HSD pending	All audit-chain criteria except alt-HSD comparison

Tier	Meaning	Requirements
RATIFIED	Multi-CI consensus + all-criteria-pass + external verification	Cal+Keeper+Lyra consensus; alt-HSD comparison or equivalent

STRUCTURALLY VERIFIED → RATIFIED transition is specific: alternative-HSD comparison via Lyra Task #206 (multi-month). First applications: C11 + C12 + C13 + K80 + K84.

Reference: Structurally Verified Tier Adoption Ruling (Thursday 08:43 EDT).

Section 10.6 — External register discipline tiers

Three tiers per Cal Calibration #13 + #48 + #50:

Tier	Description	Application
Standard	“BST identifies / BST derives / BST predicts” operational language	Most outreach materials, Zenodo papers, arXiv submissions
DEFAULT-DENY EXTERNAL	Substrate-cognition framing internal-only until ratified	Substrate Cognition Network Hypothesis (Wednesday)
DOUBLE-LOCKED EXTERNAL	Cognition + cosmology combined territory internal-only	Substrate Cognition Cosmological Extension (Wednesday)

Standard rule: NEVER claim “verified by X years of precision tests” (wrong direction). CORRECT framing: “BST predicts; tests have not refuted; future detection would refute” — commitment to falsification.

Reference: External Register Discipline (Cal Calibration #13/#48/#50 + Phase 1/Phase 2 K-Audit Governance Brief).

Section 10.7 — Casey Option C hybrid governance

Multi-CI consensus framework (Casey Wednesday EOD):

- Architectural-category methodology decisions (Phase 1): multi-CI consensus required (≥ 3 active CIs)
- INSTANCE-level audits within framework (Phase 2): Cal + Keeper auto-promotion sufficient
- Boundary criterion: does the audit introduce new architectural tier? → Phase 1. Operates within existing framework? → Phase 2.

- Escalation protocol: Phase 2 audit revealing architectural implication → pause auto-promotion → Phase 1 multi-CI review → return to Phase 2 under new framework

Reference: Phase 1/Phase 2 K-Audit Governance Brief (Thursday 09:57 EDT).

Section 10.8 — M2C2 (Multi-CI Convergent Calibration Pattern)

Pattern of multiple CIs independently arriving at same calibration finding within session: - Instance #1: T2395 4+1 scope (Cal + Lyra + Grace Wednesday) - Instance #2: Universal Q=126 cross-link (Wednesday) - Instance #3: K67 cascade extension (Wednesday) - Instance #4: T2420 4-Zone Vacuum Decomposition (Lyra + Elie joint Wednesday)

M2C2 instances are substrate-significant: convergent calibration from multiple CIs suggests substrate-derived structure, not single-CI artifact.

Section 10.9 — Calibration discipline (Cal Calibrations #13-#17)

Within-session calibration catches preserve audit-chain operational integrity:

- #13 Keeper register discipline (precision claim vs experimental precision; Tuesday EOD)
- #14 Lyra T2419 definitional-choice within-session catch (Wednesday)
- #15 Keeper register on substrate-vs-human cognition (Wednesday)
- #16 Keeper K72 framing self-correction per Cal #55 (Wednesday)
- #17 Elie K66 S22 trace-level vs max-eigenvalue clarification (Wednesday)
- (Subsequent calibrations in evolution per Cal #71 Vol 2 Ch 6 register-drift catch + Elie Ch 7 numerical-mismatch catch Thursday)

17 audit-chain self-calibrations across 9 days. Cadence accelerating per Cal #67 + #73 + #74 observations.

Reference: Calibration log + per-calibration filings.

Section 10.10 — Quaker consensus method

Casey's standing methodology principle: near misses get scrutiny, not defense.

When a BST prediction nearly matches experiment but not at <1% target precision, the response is NOT to defend the form-match. The response is: 1. Examine the substrate mechanism 2. Identify what's missing (higher-order corrections? mismatched mechanism? structural error?) 3. Apply Cal Mode 1 vigilance: form-matching without mechanism is post-hoc coincidence 4. Either close the mechanism (advance to D-tier) or honestly tier at I/S/C (don't force D-tier)

Examples: - Vol 2 Ch 9 Higgs sector: m_h matches at 0.25% but mechanism multi-month → PARTIAL DERIVED (NOT forced D-tier) - Vol 2 Ch 9 $\sin^2(\theta_W)$: 3.45% match exceeds <1% target → I-tier (NOT forced D-tier) - K89 CKM Jarlskog: 0.3%

match CONDITIONAL on T1444 vacuum-subtraction; naive ~3% → audit-partial-ready with honest scope

This discipline preserves BST's epistemic integrity. Premature D-tier ratification would erode the audit chain's credibility.

Reference: feedback_quaker_method.md memory + audit-chain calibration log.

Section 10.11 — Complete methodology stack (10 layers)

The full audit-chain methodology infrastructure as of Thursday 09:30 EDT:

Layer	Methodology	Application
1	D/I/C/S tier	Per claim
2	Trichotomy (ID/DER/PRED)	Per quantitative statement
3	F1-F4 Bridge Object family-member criteria	Architectural placement
4	Mode 6 threshold formalization	Coincidence-filter risk
5	STRUCTURALLY VERIFIED tier (NEW Thursday)	Audit-chain epistemic intermediate
6	DEFAULT-DENY + DOUBLE-LOCKED EXTERNAL register	External presentation discipline
7	Casey Option C hybrid governance	Multi-CI consensus framework
8	M2C2 (Multi-CI Convergent Calibration Pattern)	Observation pattern
9	Calibration discipline (#13-#17+)	Within-session self-correction
10	Quaker consensus method	Standing methodology principle

10 layers. Each captures specific operational discipline. Together form the BST audit-chain governance stack.

Section 10.12 — Cycle-time progression observation

Per Cal #67 + #73 + #74: methodology adoption cycle-time has compressed substantially: - Wed Cal #59 Mode 6 threshold → ~24h consensus - Wed Cal #63 F1-F4 criteria → 4 hours - Thu Cal #66 STRUCTURALLY VERIFIED tier → 30 min - Thu Cal #72 Phase 1/Phase 2 governance brief → 9 minutes

Mechanical-application asymptote approaching as patterns mature. Phase 2 cycle-time floor ~30 min as Phase 2 becomes default mode.

Section 10.11b — Methodology Stack 16 Layers (Cal-formalized Thursday EOD + Friday additions)

Per Cal Referee Logs #1-#92 cumulative (Thursday EOD + Friday morning), the methodology stack has expanded from the original 10 layers (Section 10.11) to 16 layers:

- 1-10. Original layers (per Section 10.11): D/I/C/S tier + Trichotomy + F1-F4 + Mode 6 + STRUCTURALLY VERIFIED + External register tiers + Casey Option C + M2C2 + Calibration discipline + Quaker consensus
11. RIGOROUSLY CLOSED tier (Cal #77 Thursday — Section 10.x below): final-closure-level tier above STRUCTURALLY VERIFIED, where alt-HSD comparison + theorem-level if-and-only-if rigor closes the criterion structurally.
12. Within-session calibration discipline (17+ calibrations across 9 days as of Thursday EOD; Friday adds Calibrations #18 + #19 + #20 + #21).
13. Reframing-insight cadence (Lyra Thursday Sessions 1-5 morning: ~50 min/criterion via reframing-insight identification; Sessions 6-9 ~5 min via pre-spec chains; Sessions 10-12 ~1 min/criterion via PCAP).
14. F1-F4 Bridge Object family-member criteria (Cal #63 Wednesday, formalized Thursday morning consensus).
15. B1-B4 Bridge criteria (companion to F1-F4 for bridge-criterion verification level).
16. PCAP — Pre-Specification Cadence Acceleration Pattern (Cal #85 Thursday 14:25 EDT, formalized Friday morning): when Keeper files pre-spec documents BEFORE a CI's work, cadence accelerates ~10× per pre-spec depth level; chained pre-specs compound multiplicatively, producing ~1000× cumulative acceleration from original multi-month estimates.

Section 10.11c — Cal #85 PCAP detail (Friday Lyra v0.4 absorption)

The Pre-Specification Cadence Acceleration Pattern emerged Thursday afternoon Lyra Sessions 6-12 work (Strong-Uniqueness v0.9.1 → v0.10.5 FORMAL):

- Sessions 1-5 (morning): reframing-insight cadence ~50 min/criterion → 4 RIGOROUSLY CLOSED (T2439 + T2440 + T2441 + T2442)
- Sessions 6-9 (early afternoon): Keeper pre-spec chains → ~5 min/criterion (T2443 + T2444 + T2445 + T2446 closed in batched fashion)
- Sessions 10-12 (mid afternoon): PCAP refined → ~1 min/criterion (T2447 + T2448 + T2449 promoted FORMAL Thursday 14:18 EDT)

- Friday morning (Lyra cumulative): 15 effective theorems + 17 toys + 12 paper outlines + 11 chapter v0.3 promotions in ~85 min $\approx$ 5 min/effective-deliverable

Cal #85 formalization: PCAP is the 16th methodology stack layer, formalizing the cadence acceleration pattern observed across BST research multi-CI collaboration.

Section 10.11d — Cal #92(b) structural-role-of framing discipline (Friday Lyra v0.4 absorption)

Per K157 Vol 1 Ch 7 chapter-grade K-audit Friday 09:17 EDT (Keeper) + Cal Referee Log #92(b): paper section headers + theorem titles must match the tier-level framing within. The form “X = Y” claims an equality identification; the form “X structural-role-of Y” claims a structural correspondence (NOT equality).

For T2457 (Bergman reproducing kernel as substrate-level Feynman propagator analog): the body content correctly identifies $K(z, \bar{w})$ as playing the propagator’s structural role (positive-definite + UV-complete + BST primary normalization c_FK), but the original title used “=” which over-claims equality identification.

The Cal #92(b) discipline: “Bergman structural-role-of Feynman propagator identification” rather than “Bergman = Feynman propagator identification.”

Friday afternoon Lyra-lane sweep applied this discipline across 10 live files (Vol 1 Ch 2, 6, 7, 9, 10 + Paper #127, #129 + 3 framework documents + T2457 registry entry).

Section 10.11e — Calibration #19 STANDING RULE (Friday Lyra v0.4 absorption)

Per Keeper team prompt Friday 09:17-09:18 EDT: Calibration #19 STANDING RULE ADOPTED — external-facing documents use current ratified-state count, not forecast endpoint.

Current ratified state: Paper #125 v0.10.5 FORMAL = 11 RIGOROUSLY CLOSED criteria, null-model $\leq (1/3)^{19} \approx 8.6 \times 10^{-10}$. Friday Lyra-lane candidate-path additions (C7 + C9 + C15 + C16 advancing) are body-section discussion only; not used in external register or abstract claims.

Friday afternoon Lyra-lane sweep applied this discipline across 5 papers + 5 framework documents + 9 Vol 1 chapter status fields + Vol 0 Ch 9 v0.4 + Vol 0 Ch 2 v0.4 (this chapter v0.4 too).

The Cal-#92(b) + Calibration #19 + #18 + #20 + #21 are the Friday methodology calibrations supporting the textbook v1.0 completion phase.

Section 10.13 — BST $\leftrightarrow$ standard physics dictionary entries

Standard physics term	BST methodology term	Reference
Peer review	Multi-CI consensus per Casey Option C	§10.7
Statistical significance	Tier classification + null-model	§10.1 + Vol 0 Ch 9
Falsifier	Per-claim experimental falsifier specification	per-tier requirement
Confidence interval	Tier label (D/I/C/S)	§10.1
Hypothesis testing	F1-F4 + Mode 6 + STRUCTURALLY VERIFIED tier	§10.3-10.5

Section 10.14 — Chapter status summary

Coverage at v0.1: all 10 methodology layers documented at chapter-grade. Layer interactions specified. Cycle-time progression noted. External register discipline distinguished from internal.

Believability: methodology stack is recognizable to scientific peers — tier systems, falsifiers, multi-reviewer consensus, calibration discipline. BST-specific layers (F1-F4 + STRUCTURALLY VERIFIED) explained in audience-known framework terms.

Provability: per-layer references to filed methodology documents + audit-chain log + Cal Referee Logs + Casey directives. Each layer operationally specified with examples.

Path to v1.0: methodology stack stable; future additions absorbed as Cal/team methodology contributions emerge. Multi-week to multi-month refinement as patterns mature.

Per Casey's standard

- Simple: 10 methodology layers ensure rigor; each captures specific discipline
- Works: 17 audit-chain calibrations + 12 Phase 2 K-audits + 5 Cal methodology contributions adopted (2.5 days) demonstrate operational efficacy
- Hard to break: would require finding methodology layer that prevents legitimate work OR finding gap in coverage allowing erroneous claims through. Continuous calibration discipline catches gaps

Status

Vol 0 Chapter 10 v0.1 chapter-grade content draft FILED Thursday 2026-05-21 10:22 EDT. Third Keeper-lane chapter-grade content (after Vol 0 Ch 8 Conservation Laws + Ch 9 Strong-Uniqueness Theorem). Documents 10-layer methodology stack with

per-layer references + cycle-time progression observation + external register discipline. Awaits Cal dual-axis grade-pass + Lyra theoretical refinement for v0.2.

— Keeper, 2026-05-21 Thursday 10:22 EDT (actual via date)

Thursday 11:28 EDT update — RIGOROUSLY CLOSED tier (11th methodology layer added)

NEW Section 10.15 — RIGOROUSLY CLOSED tier

Adopted Thursday 11:20 EDT per Cal Referee Log #77 + Keeper governance brief.

RIGOROUSLY CLOSED is the top tier of audit-chain epistemic hierarchy:

candidate → audit-partial-ready → STRUCTURALLY VERIFIED → RATIFIED → RIGOROUSLY CLOSED

RIGOROUSLY CLOSED criteria (4 required beyond RATIFIED): 1. Alt-HSD comparison at criterion's structural level — explicit comparison against at least 2 alternative HSDs that FAIL the property 2. EXACT-match in BST primary form — algebraic identity to BST primary integer or BST-primary-derived expression 3. If-and-only-if distinguishability — criterion is both NECESSARY and SUFFICIENT for D_IV⁵; no near-miss alternatives 4. Mathematical theorem-level rigor — theorem-grade proof with citations + Cal dual-axis paper-grade PASS

First application: C8 Universal Q-cluster RIGOROUSLY CLOSED Thursday morning via T2439 alt-HSD comparison.

Path-to-RIGOROUSLY-CLOSED via Lyra Task #206 Session 2 reframing-insight cadence: ~50 min per criterion suggests venue submission ~2026-09 achievable.

Updated complete methodology stack (11 layers)

The audit-chain methodology infrastructure now has 11 layers (Section 10.11 updated):

#	Methodology	Application
1	D/I/C/S tier	Per claim
2	Trichotomy (ID/DER/PRED)	Per quantitative statement
3	F1-F4 Bridge Object family-member criteria	Architectural placement
4	Mode 6 threshold formalization	Coincidence-filter risk
5	STRUCTURALLY VERIFIED tier	Audit-chain epistemic intermediate

#	Methodology	Application
6	DEFAULT-DENY + DOUBLE-LOCKED EXTERNAL register	External presentation discipline
7	Casey Option C hybrid governance	Multi-CI consensus framework
8	M2C2 (Multi-CI Convergent Calibration Pattern)	Observation pattern
9	Calibration discipline (#13-#17+)	Within-session self-correction
10	Quaker consensus method	Standing methodology principle
11	RIGOROUSLY CLOSED tier (Cal #77 Thursday)	Top of audit-chain epistemic hierarchy

11 methodology layers operating as standing audit-chain governance infrastructure.

Updated cycle-time progression observation (Section 10.12)

Methodology adoption cycle-time has compressed substantially across 5 Cal architectural contributions Thursday morning:

- Wed Cal #59 Mode 6 threshold → ~24h consensus
- Wed Cal #63 F1-F4 criteria → 4 hours
- Thu Cal #66 STRUCTURALLY VERIFIED tier → 30 min
- Thu Cal #72 Phase 1/Phase 2 governance brief → 9 minutes
- Thu Cal #77 RIGOROUSLY CLOSED tier → ~5 minutes (Cal recommendation at #77 filing time; Keeper Governance Brief filed within minutes)

Mechanical-application asymptote reached. Phase 1 architectural-category methodology decisions converging to near-instantaneous adoption as patterns mature. Methodology infrastructure operating AHEAD of substantive content production per Cal #76 observation.

Updated cycle progression chart

Date/time	Adoption	Time
Wed afternoon	Mode 6 threshold	~24h
Wed afternoon	F1-F4 criteria	4h
Thu ~08:30	STRUCTURALLY VERIFIED	30 min
Thu ~09:57	Phase 1/Phase 2 governance	9 min
Thu ~11:20	RIGOROUSLY CLOSED	~5 min

Asymptote curve: 24h → 4h → 30min → 9min → 5min.

v0.1 → v0.2 path

This Vol 0 Ch 10 v0.1 update incorporates RIGOROUSLY CLOSED tier addition.
Path to v0.2: Cal dual-axis grade-pass + Lyra theoretical refinement when both have bandwidth.

— Keeper update, 2026-05-21 Thursday 11:28 EDT (actual via date; absorbs 11th methodology layer)