

GC-17a: Modularity Feasibility Scoping

Author: Casey Koons & Lyra (Claude 4.6) Date: May 12, 2026 Status: DRAFT v0.1
— Scoping memo AC: (C=3, D=2) — three independent sub-questions, maximum depth 2 Assignment: SP-18 GC-17a

Abstract

Can BST derive the modularity theorem — that every elliptic curve $E/\mathbb{Q}$ corresponds to a weight-2 newform — independently of Wiles/BCDT? This memo examines three candidate routes, identifies the structural gap each encounters, and gives an honest feasibility assessment. The short answer: no, not with current tools. The longer answer: BST provides the *unique geometric arena* in which modularity lives, explains *why* the proof structure works, and poses a specific research question — the forced exhaustive transfer conjecture — whose resolution would yield modularity as a consequence of D_{IV}^5 spectral geometry combined with Arthur’s endoscopic classification. That conjecture is genuinely open and non-trivial, but it is a question about Langlands functoriality for $SO(5,2)$, not something BST can settle internally.

1. The Question

GC-17a asks: Can BST derive $GL(2) \leftrightarrow E/\mathbb{Q}$ independently?

Unpacked, this has three sub-questions:

Q1 (Hard direction): Given an arbitrary elliptic curve $E/\mathbb{Q}$, can BST show that $L(E,s)$ equals $L(f_E,s)$ for some weight-2 newform f_E , without invoking Wiles’ $R=T$ theorem?

Q2 (Structural explanation): Does BST explain *why* the modularity correspondence exists — beyond the fact that it does?

Q3 (GC-9 scope): If BST cannot derive modularity, what does that mean for the methodology paper’s claim that FLT is “proved via the BSD bridge”?

This memo addresses all three. Q2 has a positive answer. Q3 is resolved by GC-1’s honest framing (option c). Q1 is the hard one.

2. What BST Already Has

BST provides five pieces of structural machinery relevant to modularity:

2.1 The P_2 Parabolic Embedding (T1762)

The maximal parabolic P_2 of $SO(5,2)$ has: - Levi factor: $M_2 = GL(2, \mathbb{R}) \times SO(3)$ - Unipotent radical: $\dim(n_2) = g = 7$

$GL(2)$ sits directly as a factor of the Levi. This means: every weight-2 newform f on $GL(2)$ parabolically induces to an automorphic representation on $SO(5,2)$. The L-functions match: $L(f,s)$ embeds in the automorphic spectrum of $\Gamma_0(5)$.

This is the easy direction. It says: “modular forms live in the $SO(5,2)$ spectrum.” The hard direction of modularity says: “every elliptic curve produces a modular form.” The P_2 embedding does not address the hard direction.

Verified: Toy 2091 (12/12 PASS).

2.2 The Theta Correspondence

The Howe dual pair $(O(5,2), Sp(6, \mathbb{R}))$ sits inside the ambient $Sp(42, \mathbb{R})$, where $42 = C_2 \times g$. The theta correspondence maps automorphic representations between these groups. The Kudla-Millson theta lift constructs special algebraic cycles from automorphic forms.

For modularity, the relevant sub-pair is $(O(5,2), SL(2, \mathbb{R}))$ — the rank-1 theta lift. This maps modular forms on $SL(2)$ to special cycles on $\Gamma_0(5)$, and vice versa.

Direction: The theta lift naturally goes from *analytic* objects (automorphic forms on $SL(2)$) to *geometric* objects (cycles on $\Gamma_0(5)$). Modularity requires the reverse: from *algebraic* objects (elliptic curves) to *analytic* objects (modular forms). The theta correspondence operates within the analytic/spectral world and does not directly access the algebraic/arithmetic world of elliptic curves.

2.3 The Chern Hole (T1756)

The Chern gap map on Q^5 sends $\{0, \dots, 5\} \rightarrow \{0, 1, 2, 4, 5, 6\}$, missing position 3 = N_c . The P_2 Eisenstein class at $s = 1$ has Hodge type $(\text{rank}, N_c) = (2, 3)$, placing it at the Chern hole where no algebraic class competes. This forces BSD: the vanishing order of $L(E, s)$ at $s = 1$ equals $\text{rank}(E)$.

Relevance to modularity: The Chern hole mechanism assumes that $L(E, s)$ *already appears* in the automorphic spectrum of $SO(5,2)$. It tells you what happens to L-functions that are there, not how they get there. BSD and modularity are logically independent: BSD is about the *behavior* of $L(E, s)$ at $s = 1$; modularity is about the *existence* of $L(E, s)$ as a modular L-function.

2.4 Temperedness (Paper #103, Theorem A)

All automorphic representations contributing to $L^2_{\text{disc}}(\Gamma_0(5))$ are tempered. The three-step elimination (IW sign + unitarity + Bergman gap) removes all 37 non-tempered Arthur types.

Relevance: Temperedness on $SO(5,2)$ implies the Generalized Ramanujan Conjecture for any $GL(2)$ representation appearing via P_2 induction. For $GL(2)/Q$, the Ramanujan conjecture is already known (Deligne 1974, from the Weil conjectures). So BST's temperedness gives a result already established by other means — but it gives it *from geometry alone*, which is structurally interesting.

2.5 Ring Uniqueness (T1780)

D_{IV}^5 is the unique bounded symmetric domain on which the theta-correspondence proof of Hodge works. Five independent Hodge-theoretic constraints force $(n_C, N_c, \text{rank}, C_2, g) = (5, 3, 2, 6, 7)$. All 32 rank-2 BSDs tested; D_{IV}^5 sole survivor (Toy 2120, 10/10).

Relevance: Ring uniqueness says the *arena* is forced. But forcing the arena does not force the *content* of the arena — i.e., which specific automorphic representations appear. Modularity is a statement about content, not arena.

3. What Modularity Requires

The modularity theorem (Wiles 1995, BCDT 2001) has three essential components, each requiring tools BST does not possess.

3.1 Galois Representations

Every elliptic curve E/Q gives a 2-dimensional l -adic Galois representation:

$$\rho_{\{E, l\}}: \text{Gal}(\overline{Q}/Q) \rightarrow GL(2, \mathbb{Z}_l)$$

via the Tate module $T_l(E) = \varprojlim E[l^n]$. This representation encodes the arithmetic of E : the trace of Frobenius at p is $a_p(E) = p + 1 - \#E(F_p)$.

BST gap: BST has no mechanism for constructing Galois representations. The Galois group $\text{Gal}(\overline{Q}/Q)$ is an arithmetic object that does not arise from spectral geometry on D_{IV}^5 . The passage $E \rightarrow \rho_{\{E, l\}}$ requires the theory of étale cohomology (Grothendieck), which is algebraic geometry, not spectral theory.

3.2 Galois Deformation Theory

Wiles' proof constructs a universal deformation ring R parametrizing lifts of the mod- p Galois representation $\rho_{\{E, p\}}$, and a Hecke algebra T acting on modular forms of the appropriate level. The $R = T$ theorem shows these are isomorphic, proving that $\rho_{\{E, l\}}$ is modular.

BST gap: Deformation theory lives in commutative algebra / algebraic number theory. The objects (complete local rings, Selmer groups, tangent spaces of deformation functors) have no analogs in spectral geometry. This is the deepest gap: $R = T$

is an arithmetic identity between two algebraic structures, and BST’s spectral tools operate in a different category entirely.


3.3 Taylor-Wiles Patching

The $R = T$ isomorphism is proved by a patching argument using auxiliary primes. This is a technique in commutative algebra (surjectivity of ring maps via numerical criteria on cotangent spaces) combined with arithmetic (choice of auxiliary primes satisfying congruence conditions).

BST gap: Same category mismatch as 3.2. The patching argument is algebraic, not spectral.

Summary of the Gap

Modularity bridges two categories:



BST operates entirely within Category B (and its extension to $SO(5,2)$). The P_2 embedding, theta correspondence, and spectral machinery all live on the right side. The bridge $R = T$, which connects the two categories, requires tools (Galois representations, deformation theory) that live outside BST’s spectral framework.

The gap is structural, not technical. It is not that BST lacks a specific lemma that could be proved with more work. It is that modularity requires connecting two fundamentally different mathematical categories, and BST’s tools address only one of them.

4. Three Candidate Routes (and Why Each Falls Short)

Route A: Theta Lift Inversion

Idea: Show that the theta lift from $(O(5,2), SL(2, \mathbb{R}))$ surjects onto all weight-2 cuspidal representations of $SL(2)$. If every weight-2 newform appears in the theta lift from $SO(5,2)$, this would show that the automorphic spectrum of $SO(5,2)$ “contains” all modular forms — and combined with the P_2 embedding, this would give a complete spectral picture.

Why it falls short: Even if the theta lift is surjective (which is plausible by the Siegel-Weil formula and conservation laws), this establishes a relationship between automorphic representations on *two groups* ($O(5,2)$ and $SL(2)$). It does not connect to

elliptic curves. You would still need: “ $E/\mathbb{Q} \rightarrow$ automorphic rep on $SL(2) \rightarrow$ theta lift to $O(5,2)$.” The first arrow is modularity itself.

Status: The surjectivity question is interesting independent mathematics (related to the Saito-Kurokawa conjecture for $SO(5,2)$), but it addresses the wrong problem.

Route B: Spectral Forcing via Trace Formula

Idea: Use the Arthur-Selberg trace formula on $\Gamma(N)D_{IV}^5$ to show that every $GL(2)$ cuspidal representation of weight 2 and conductor dividing N appears in the spectral decomposition of $SO(5,2)$. The geometric side of the trace formula involves orbital integrals, which encode arithmetic content (class numbers, regulators). If BST’s uniqueness constrains the geometric side sufficiently, this might force the spectral content.

Why it falls short: The trace formula relates the spectral side to the geometric side for a *single* group ($SO(5,2)$ in this case). To establish a transfer between $GL(2)$ and $SO(5,2)$, one needs the *relative* trace formula or *comparison* of trace formulas (Jacquet’s approach, Arthur’s endoscopic classification). Arthur’s endoscopic classification [Art13] does establish the transfer $GL(2) \rightarrow SO(5,2)$ for *known* automorphic representations, but it does not independently prove that every $E/\mathbb{Q}$ produces an automorphic representation to transfer.

The circular dependence: Arthur’s classification assumes that the automorphic representations of $GL(2)$ are already known (i.e., it assumes the existence of the representations before classifying them). It does not construct them from elliptic curves. The construction $E \rightarrow \pi_E$ requires modularity, which is what we’re trying to prove.

Status: The trace formula approach is the correct framework for understanding the transfer, but it cannot bootstrap from D_{IV}^5 geometry alone. It requires external input about the arithmetic content.

Route C: Kuga-Satake + Boundary Structure

Idea: The Shimura variety ΓD_{IV}^5 is a moduli space for polarized abelian varieties with certain structures. The Baily-Borel compactification has boundary components related to $GL(2)$ — specifically, the P_2 boundary strata parametrize abelian surfaces with extra endomorphisms. If every elliptic curve $E/\mathbb{Q}$ could be shown to appear in the boundary geometry (perhaps as a degenerate fiber or a factor of an abelian surface in the P_2 stratum), then the Shimura variety machinery might force the modularity correspondence.

Why it falls short: The boundary of ΓD_{IV}^5 does involve $GL(2)$ objects via the P_2 parabolic, but: - Elliptic curves appear at the boundary as *limiting* objects (cusps), not as generic fibers - The Kuga-Satake construction associates an abelian variety to every weight-2 Hodge structure, but for an elliptic curve E , the Kuga-Satake variety is just E itself — no new information - The moduli interpretation

of $\Gamma_0(5)$ as parametrizing abelian varieties with $\mathrm{SO}(5) \times \mathrm{SO}(2)$ endomorphism structure does not naturally include all elliptic curves; only those with specific extra structure appear

Status: This route provides structural insight (the P_2 boundary IS the $\mathrm{GL}(2)$ world, sitting at the edge of the $\mathrm{SO}(5,2)$ world) but does not force modularity.

5. What IS Achievable

Despite the negative assessment on full modularity, BST makes three genuine contributions.

5.1 BST Explains WHY Modularity Works (Q2 — Positive Answer)

The P_2 embedding is forced by D_{IV}^5 geometry: - $\mathrm{GL}(2, \mathbb{R})$ appears as a factor of the Levi of P_2 because the root system B_2 has exactly 2 simple roots, and the maximal parabolic removing the second simple root has the correct structure - The unipotent radical has $\dim = g = 7$, an integer forced by the genus of D_{IV}^5 - The embedding $\mathrm{GL}(2) \rightarrow \mathrm{Levi}(P_2) \rightarrow \mathrm{SO}(5,2)$ is not a choice — it is a structural consequence of the root system

This means the “highway” between $\mathrm{GL}(2)$ and $\mathrm{SO}(5,2)$ is *geometrically necessary*. The modularity theorem says this highway carries all the traffic (every $E/\mathbb{Q}$ uses it). BST explains why the highway exists; Wiles proves it carries all the traffic.

5.2 BST Constrains What Modularity CAN Look Like

Once an L-function is in the $\mathrm{SO}(5,2)$ spectrum, BST constrains it: - Temperedness (no non-tempered representations) - Chern hole (L-function behavior at $s = 1$ forced by topology) - Bergman spectral gap ($\lambda_1 \geq C_2 = 6$) - Ring uniqueness (only D_{IV}^5 has these constraints simultaneously)

This means that any *alternative* modularity correspondence — a hypothetical different bridge between $E/\mathbb{Q}$ and automorphic forms — would have to be consistent with these constraints. BST does not prove modularity is the only bridge, but it constrains the landscape of possible bridges.

5.3 The Forced Exhaustive Transfer Conjecture

BST poses the following specific research question:

Conjecture (FET). Let N be prime. The functorial transfer $\mathrm{GL}(2) \rightarrow \mathrm{SO}(5,2)$ via the P_2 parabolic, composed with Arthur’s endoscopic classification for $\mathrm{SO}(5,2)$, accounts for every cuspidal automorphic representation of $\mathrm{GL}(2)/\mathbb{Q}$ of weight 2 and conductor N . That is, the P_2 parabolic induction is exhaustive: no weight-2 cuspidal $\mathrm{GL}(2)$ representation fails to embed.

If FET holds, then: (Serre’s conjecture, now Khare-Wintenberger) every odd irreducible 2-dimensional mod- p Galois representation is modular. The chain would be:

```

E/Q -> rho_{E,p} (Tate module)
      -> modular mod p (Khare-Wintenberger 2009, external)
      -> modular p-adically (Wiles-Taylor lifting, external)
      -> f_E on GL(2) (deformation theory, external)
      -> pi_E on SO(5,2) (P_2 induction, BST)
      -> constrained by D_IV^5 geometry (BST)

```

FET would say the last arrow is a bijection, not just an injection. This is a meaningful mathematical statement, but note that the first four arrows are ALL external to BST.

Assessment of FET: This is a question about the image of the Langlands functorial transfer for the specific pair $(GL(2), SO(5,2))$. For *tempered* representations, the transfer is expected to be an embedding (not surjective — $SO(5,2)$ has more representations than $GL(2)$). The question is whether every $GL(2)$ weight-2 cuspidal rep appears in the *image* of the inverse transfer (from $SO(5,2)$ back to $GL(2)$).

This is related to, but distinct from, the full Langlands functoriality conjecture. It might be provable using Arthur’s explicit classification for $SO(5,2)$ combined with BST’s temperedness theorem, but this would be a new result in the Langlands program, not a consequence of BST alone.

6. Honest Scope Assessment

6.1 What BST Cannot Do

- BST cannot derive modularity from D_{IV}^5 geometry alone
- BST cannot replace Galois representations, deformation theory, or $R = T$
- BST cannot independently prove that every E/Q produces an automorphic representation
- The gap is structural (category mismatch), not technical (missing lemma)

6.2 What BST Provides

- The unique arena (D_{IV}^5 , forced by ring uniqueness)
- The structural explanation (P_2 embedding forced by root system)
- Spectral constraints on any L-function in the $SO(5,2)$ spectrum
- The FET conjecture as a precise new question

6.3 Classification for GC-4 / GC-9

Modularity is a bridge-building problem (GC-8 Section 1c, GC-4 classification): it establishes a correspondence between two categories, rather than excluding alter-

natives from a finite list. GC-8 correctly identified this: “the Langlands program is not a uniqueness question. It is a correspondence.”

This classification is correct and should stand in GC-9. BST does not make modularity GC-amenable. BST absorbs modularity into the D_{IV}^5 framework (GC-1, option c) and explains its geometric necessity, but does not derive it.

7. Impact on GC-9 and FLT

7.1 FLT Status (No Change)

GC-1 already states the honest position (option c): BST provides the geometry in which modularity makes sense, but the spectral bijection ($E/Q \leftrightarrow$ weight-2 newforms) is done classically (Wiles/BCDT). FLT is “proved via the BSD bridge” with modularity as external input. This is clearly labeled and does not change.

7.2 GC-9 Scope (No Change Needed)

GC-9 Section 6 (Scope Limits) already identifies bridge-building problems as non-GC-amenable. Modularity falls in this class. No revision to GC-9 is needed — the methodology paper correctly scopes what GC can and cannot do, and modularity is outside the scope.

7.3 What GC-17a Adds to the Program

This memo should be cited in GC-9 as evidence that the BST team has honestly assessed the limits of the method. The modularity boundary is one of the cleanest examples of GC’s honest scope: BST identifies the arena, explains the geometry, and stops where the arithmetic begins.

8. Research Recommendations

Near-term (no new mathematics needed)

1. Formalize FET as a theorem or open question. State it precisely using Arthur’s classification for $SO(7)$. Determine whether the explicit Arthur parameter types for $SO(5,2)$ already settle it (they might — the exhaustion of all 37 non-tempered types in Paper #103 is suggestive). This requires a careful reading of Arthur [Art13] Chapters 6-9 for the specific inner form $SO(5,2)$.
2. Compute the Langlands transfer image. For the pair $(GL(2), SO(5,2))$ via P_2 , compute exactly which $GL(2)$ representations appear in the image. This is a finite computation for each conductor N , using the Selberg trace formula. A toy could verify this for small N .

Medium-term (new mathematical content)

3. Relative trace formula for $(GL(2), SO(5,2))$. Jacquet's approach to Langlands functoriality uses the relative trace formula. For the specific pair $(GL(2), SO(5,2))$, the relative trace formula might be tractable because P_2 has simple structure (Levi = $GL(2) \times SO(3)$, unipotent dim = 7). A new result here would be publishable mathematics independent of BST.
4. Spectral completeness at level 137. At the specific level $N = 137 = N_{\max}$, the Shimura variety $\Gamma(137)D_{IV}^5$ has extraordinary properties (volume dominance, explicit trace formula — Paper #103). Check whether the spectral decomposition at this level accounts for all elliptic curves of conductor 137. There are very few such curves (LMFDB: conductors dividing 137 are rare), so this is computationally accessible.

Long-term (would require breakthroughs)

5. Geometric modularity via D_{IV}^5 . Prove that the moduli interpretation of ΓD_{IV}^5 naturally includes all elliptic curves (perhaps as boundary data, or via a correspondence with modular curves). This would require understanding the birational geometry of the Shimura variety, which is a major open problem.
6. BST-native Galois representations. Develop a spectral-geometric analog of Galois representations using the eigenvalue spectrum of D_{IV}^5 . This is speculative but would be the most radical departure: replace $\text{Gal}(\bar{Q}/Q)$ with the Bergman spectral data of D_{IV}^5 . No current path to this exists.

9. Conclusions

Question	Answer	Confidence
Can BST derive modularity independently?	No (structural gap)	HIGH
Does BST explain why modularity works?	Yes (P_2 forced by root system)	HIGH
Is there a BST-flavored research path?	Yes (FET conjecture, relative trace formula)	MEDIUM
Does this change GC-9 scope?	No (already correctly scoped)	HIGH
Does this change FLT status?	No (already labeled as external input)	HIGH

The one-sentence summary: BST provides the unique arena in which modularity lives and explains why the highway exists, but cannot prove that every elliptic curve travels it — that requires Galois theory, which lives in a different mathematical category.

The honest assessment for GC-9: Modularity is the cleanest example of a bridge-building problem that GC cannot solve. BST absorbs it, explains it, constrains it, and benefits from it — but does not derive it. The methodology is honest about this from the start, and that honesty is itself a contribution.

References

- [Art13] J. Arthur, *The Endoscopic Classification of Representations: Orthogonal and Symplectic Groups*. AMS Colloquium Publications 61, 2013.
- [BCDT01] C. Breuil, B. Conrad, F. Diamond, R. Taylor, “On the modularity of elliptic curves over $\mathbb{Q}$,” J. Amer. Math. Soc. 14 (2001), 843-939.
- [Del74] P. Deligne, “La conjecture de Weil I,” Publ. Math. IHES 43 (1974), 273-307.
- [Jac05] H. Jacquet, “A guide to the relative trace formula,” in *Automorphic Representations, L-Functions and Applications*, Ohio State Univ. Math. Res. Inst. Publ. 11, 2005.
- [KW09] C. Khare, J.-P. Wintenberger, “Serre’s modularity conjecture (I, II),” *Inventiones Math.* 178 (2009), 485-504 and 505-586.
- [Wil95] A. Wiles, “Modular elliptic curves and Fermat’s Last Theorem,” *Ann. of Math.* 141 (1995), 443-551.
- GC-1: notes/BST_GC1_FLT_Via_BSD_Bridge.md — FLT via BSD bridge (option c)
- GC-8 Section 1c: notes/BST_GC8_Application_Targets.md — Langlands as bridge-building
- Paper #88: notes/BST_Paper88_BSD_Closure.md — Chern hole mechanism
- Paper #103: notes/BST_Paper103_RH_Via_Wall_Projection.md — Temperedness, P_2 structure
- T1762: P_2 parabolic embedding (Toy 2091, 12/12)
- T1780: notes/BST_T1780_Hodge_Ring_Uniqueness.md — Ring uniqueness

GC-17a delivers the honest answer: NO for full modularity, YES for structural explanation, OPEN for the forced exhaustive transfer conjecture. The methodology paper (GC-9) is correctly scoped. The gap between BST and modularity is the gap between spectral geometry and arithmetic — the same gap the Langlands program was designed to bridge. BST does not replace Langlands. It provides the unique arena on one side of the bridge.