

Grace — A1: the radial D_H non-CPL discriminator (PRE-REGISTERED) (2026-08-08, K1291 Lane 2)

Lane-2 (Grace [?] Elie): ready the radial D_H discriminator — the model-independent test of whether DESI's phantom crossing is REAL or a CPL-parametrization artifact. BST predicts $w(a)$ relaxing to -1 FROM ABOVE ($w_a > 0$, completely-monotone, NO crossing — F799 double-derived). DESI's crossing rides the CPL linear form which ALLOWS crossing by construction. Target-innocent: BST shape NEVER fit to DESI; pre-registered before the merge. Λ Structural; only the shape bankable.

The observable (why radial D_H is the right probe)

$D_H(z) = c/H(z)$ — the RADIAL Hubble distance, measured directly per-redshift-bin by radial BAO, without assuming any $w(a)$ form. $H(z)^2 = H_0^2[\Omega_m(1+z)^3 + \Omega_{DE} \cdot f_{DE}(a)]$, $f_{DE}(a) = \exp(3 \int_{a^1}^1 (1+w)/a' da')$. Unlike the CPL fit (which imposes a linear $w(a)$ and thereby *permits* a crossing), the binned $D_H(z)$ is parametrization-independent — so it can decide whether the crossing is in the *data* or in the *parametrization*.

The discriminator (the signature, pre-registered)

Take the residual $R(z) = D_H(z)/D_{H,\Lambda CDM}(z) - 1$ across the DESI range ($z \approx 0.3-2.5$): - CPL with phantom crossing (DESI-like $w_0 = -0.75$, $w_a = -0.80$): $R(z)$ FLIPS SIGN — negative at low z (-3% at $z=0.3$), through zero near $z \approx 1.3$, positive at high z ($+0.7\%$ at $z=2.5$). The sign-flip IS the phantom-crossing fingerprint ($w < -1$ in the past pushes f_{DE} the other way). - BST (completely-monotone, no crossing): $R(z)$ is MONOTONE, one sign (all-negative here, $-4.4\% \rightarrow -1.7\%$), NO flip. A completely-monotone $w(a) > -1$ (F799) cannot produce a residual sign-flip — the no-crossing property forbids it structurally.

★ PRE-REGISTERED PASS/FAIL (target-innocent; decided on the model-independent $D_H(z)$, not the CPL fit)

model-independent binned radial $D_H(z)$ shows...	verdict
a monotone $R(z)$ (one sign, no flip)	the phantom crossing was a CPL artifact $\rightarrow$ BST's no-crossing SURVIVES (a cleaner fit, no pathology)
a genuine sign-flip in $R(z)$ near $z \approx 1-1.5$ that survives model-independent reconstruction	the crossing is REAL $\rightarrow$ BST's bleed mechanism is genuinely falsified (K1040's own kill condition fires)

Both outcomes pre-committed. The test is the *monotonicity of the radial residual*, NOT goodness-of-fit to a chosen parametrization.

Structural claim (target-innocent, to confirm at merge)

The discriminator's power is that complete-monotonicity (F799) [?] no residual sign-flip for ANY BST shape — it does not depend on the exact curve. So the qualitative prediction (monotone $R(z)$) is forced by F799, independent of the amplitude. The *exact* $R(z)$ (and where it's largest — here $z \approx 0.5$, $R \approx -5\%$) rides Elie's exact $w(a)$ (F797 spectral gap λ_1 = the rate; F799 = the completely-monotone shape).

Merge (Elie + Grace)

- Elie's half: the exact BST $w(a)$ — λ_1 (rate, F797) + completely-monotone shape (F799), verify $w_a > 0$ / no crossing. Output: BST's exact $w(a)$ curve.
- My half (done): the $D_H(z)$ machinery + the pre-registered sign-flip discriminator (above). At merge: plug Elie's exact $w(a)$ into $D_H(z)$, overlay the model-independent DESI radial-BAO $D_H(z)$, read the residual monotonicity.

- Then: the merged curve is laid against the discriminator; the sign-flip decides. We are on the “wrong side” of the CPL headline (known, pre-registered) — the question is whether that headline survives the non-CPL radial reconstruction.

Discipline

Target-innocent throughout — BST’s shape is a representative monotone form here ($d_0=0.25$), NEVER fit to DESI; the CPL curve uses DESI’s published (w_0, w_a) only to SHOW the phantom signature, not to tune BST. Λ stays Structural; the $w(a)$ SHAPE is Structure-Derived (F799); the DESI comparison is a pre-registered falsifier, not a fit. Machinery reusable ($D_H(z)$ for arbitrary $w(a)$). Nothing pushed.

MERGE (Elie + Grace, K1294) — structural verdict IN; data-verdict PENDING the real binned table

Elie’s exact rate landed: $\lambda_1 = C_2 = 6$ (Helgason spectral gap; power-law, not exponential). *[CORRECTED 2026-08-08: the parenthetical “Obata: = the Einstein constant” is RETRACTED — Obata is sphere-only; the honest relation is at most $\lambda_1 = \text{rank} \times \text{Einstein-constant}$ (Kähler–Einstein), gated on the operator ID. The rate = $C_2 = 6$ as a number stands; the Einstein-constant identification does not.]* Ran BST’s $w(a) = -1 + A \cdot a^6$ (rate-6 monotone) through the D_H machinery vs the DESI-DR2 CPL best-fit ($w_0=-0.838$, $w_a=-0.62$, K1120-verified):

z_{eff}	0.30	0.50	0.71	0.93	1.32	1.49	2.33	shape
BST ($A=0.16$, rate 6)	−1.61	−1.47	−1.21	−0.95	−0.62	−0.52	−0.24	MONOTONE (every A)
CPL DESI-fit	−1.74	−1.27	−0.55	+0.08	+0.71	+0.84	+0.85	SIGN- FLIP (phan- tom)

Structural verdict (target-innocent, amplitude-independent): BST’s rate-6 completely-monotone $w(a)$ gives a monotone radial residual for every amplitude A ; the CPL phantom fit flips sign near $z \approx 0.9$. Cleanly distinguishable in the radial D_H shape. (Amplitude A left free — NOT fit to DESI; the BST w_0 is unresolved-between-two-forms per Cal §164, and the discriminator does not depend on it.)

★ DATA-VERDICT PENDING (honest — do NOT over-claim): the definitive read — does the MODEL-INDEPENDENT binned radial-BAO $D_H(z)$ show monotone or flip? — requires the real DESI DR2 BAO D_H/r_d table + covariance (each z -bin’s $H(z)$, no $w(a)$ assumed). I attempted to verify at source (WebFetch arXiv:2503.14738 abstract + HTML) — the binned table did not extract cleanly, and I refuse to substitute remembered numbers or use the CPL fit as data (circular). So: the discriminator is BUILT and structurally decisive; the verdict awaits the proper data pull (DESI DR2 BAO table + covariance → compute $R(z)$ + its monotonicity significance). That pull is the one remaining concrete step.

Honest status: BST predicts a monotone radial residual (no crossing), forced by complete-monotonicity (F799/Bernstein) + rate $C_2=6$ (the spectral gap; NOT “Obata” — retracted). DESI’s CPL *headline* prefers a crossing. Whether the *data* (beyond CPL) requires it is the pre-registered, both-ways test — still open, pending the binned reconstruction. Λ Structural; shape Structure-Derived; never fit to DESI; never “BST derives dark energy.”

★ DATA VERDICT (K1295) — real DESI DR2 radial BAO pulled at source; the crossing is NOT in the radial data

Got the real table (ar5iv Table 4, arXiv:2503.14738, verified at source — $r_d=147.05$; D_H/r_d for LRG1/LRG2/LRG3+ELG1/ELG2/ with errors + D_M-D_H correlations).

compressed CMB anchor (single distance, not the full Planck likelihood); the definitive significance is DESI's full analysis, not my fit. Λ Structural; nothing fit to force a match; never "BST derives dark energy."

Net across both: radial BAO alone $\rightarrow$ BST survives clean ($\chi^2/\text{dof}=1.0$, no crossing in the direct $H(z)$); BAO+CMB $\rightarrow$ BST shares ΛCDM 's $\sim 3\sigma$ tension with the dynamical-DE preference. Both $<5\sigma$ and contested. Honest, calibrated both ways.

★ RE-FIRED on Keeper's AUTHORITATIVE table (K1436, Cobaya ALL_GCcomb) — verdict CONFIRMED (2026-08-12)

Keeper (K1436) pulled the machine-readable DESI DR2 BAO D_H/r_d + covariance from the Cobaya pipeline files (NOT the ar5iv HTML I used in K1295, NOT the CPL fit). Block-diagonal covariance $\rightarrow$ the 6 D_H points are mutually independent $\rightarrow$ diagonal errors, clean.

2-param ΛCDM fit to the 6 verified D_H/r_d points: $\Omega_m=0.270$, $\chi^2=3.80$, $\chi^2/\text{dof} = 0.95$ (excellent). Residuals (data- ΛCDM)/ σ : $-0.54, -0.98, +1.31, +0.57, -0.16, -0.69$ ($z = 0.51 \rightarrow 2.33$). - Pattern $--++--$ = a mid- z bump (max 1.3σ), NOT the phantom signature (a single MONOTONE $\rightarrow +$ crossing near $z \approx 0.9$). - Monotone-trend test: $\text{corr}(\text{residual}, z) = -0.085 \approx 0$. Phantom needs $\text{corr} \approx +1$; high- z residuals are NEGATIVE (opposite of phantom's $+$ at high z). - $\nexists$ the DESI phantom crossing is NOT in the model-independent radial $D_H(z)$. BST's monotone/no-crossing (F799) SURVIVES the radial-BAO probe — now on authoritative Cobaya data with proper covariance.

Scope (unchanged, honest): radial-BAO only. Full BAO+CMB(+SNe) still leans dynamical-DE $\sim 3\sigma$ (BST shares ΛCDM 's tension; K1297); the phantom is SNe+combination-driven, not the direct $H(z)$. Λ Structural; shape Structure-Derived; nothing fit to DESI. The definitive radial verdict is now on source-verified data — BST not falsified, crossing absent from the most direct $H(z)$ measurement.

★ SNe-SIDE FEASIBILITY (2026-08-13, toward pred_125's teeth) — model-level; NO SNe data yet

The phantom preference is SNe-driven. Q: can BST-monotone $H(z)$ (no crossing) + a SNe calibration offset reconcile the combined fit? Computed the SNe residual the DESI CPL-phantom fit ($w_0=-0.838$, $w_a=-0.62$) absorbs over ΛCDM (= BST-monotone, tilt $\rightarrow 0$ as the radial data prefer), $\Delta\mu = \mu_{\Lambda\text{CDM}} - \mu_{\text{phantom}}$, constant M marginalized:

z :	0.05	0.10	0.20	0.30	0.50	0.70	1.00	
$\Delta\mu$:	-13.8	-7.4	+1.3	+6.0	+8.2	+5.8	-0.1	mmag

- SIZE: ~ 0.022 mag peak-to-peak — at the LOW end of the known SNe calibration-systematic budget (~ 0.02 - 0.05 mag; Pantheon+ low- z -anchor / DES-SN5YR low- z -vs-Hubble-flow debate). The phantom pull is SMALL $\rightarrow$ absorbable by a systematic.
- SHAPE: a hump DOMINATED BY THE LOW- z POINTS (-13.8 mmag at $z=0.05$, largest deviation; shrinking by $z \approx 0.2$). This is exactly the signature of a low- z anchor offset — the residual is worst precisely where the SNe calibration debate lives ($z < 0.1$). NOT a random shape; it's low- z -driven.
- VERDICT (feasibility, honest): BST-monotone $H(z)$ + a ~ 0.02 mag low- z SNe calibration offset PLAUSIBLY reconciles the $\sim 3\sigma$ combined phantom preference. The pull is small and low- z -dominated — consistent with, not proof of, a systematic. The full teeth-earning needs the ACTUAL Pantheon+/DES-SN5YR residuals fit jointly with BST-monotone $H(z)$ + a low- z -offset nuisance parameter (does it kill the $\sim 3\sigma$?). Keeper chasing the SNe data. Λ Structural; nothing fit to force a match.

★★ REAL SNe FIT (2026-08-13, K1442 DES-SN5YR, 1820 SNe) — pred_125 teeth EARNED at the diagonal-covariance level

Fit three models to the DES-SN5YR distance vector (diagonal stat+sys errors; marginalizing absolute mag M): | model | params (over ΛCDM) | χ^2 | χ^2/dof | |—|—|—|—| | ΛCDM | — | 1516.9 | 0.834 | | CPL phantom (w_0, w_a) | +2 | 1510.3 | 0.832 | | BST-monotone + low- z offset | +2 (A, off) | 1507.1 | 0.830 | - FAIR comparison: CPL and

BST+offset BOTH have 4 params (2 extra over Λ CDM). Equal footing. - RESULT: BST-monotone+offset best-fits to $A=0$ (NO tilt $\rightarrow \Lambda$ CDM) + a +0.05 mag low- z calibration offset, and fits the SNe BETTER than the phantom ($\Delta\chi^2 = 3.2$ in BST+offset's favor). So a NO-CROSSING model + a low- z offset of the known magnitude (0.02-0.05 mag, the Pantheon+/DES-SN5YR anchor debate) fits DES-SN5YR at least as well as the phantom crossing. - $\square$ the phantom crossing is NOT required by the SNe — it is degenerate with a low- z calibration offset. This EARNS pred_125's teeth at the diagonal level: the crossing lives in the calibration, not the physics. - CAVEATS (load-bearing, honest): (1) DIAGONAL stat+sys only — the full STAT+SYS covariance (6.2MB .npz, correlated systematics) is the DEFINITIVE test; it gives systematics MORE freedom, which would only STRENGTHEN the degeneracy (BST+offset). numpy broken locally $\rightarrow$ need the covariance in a loadable form (Keeper). (2) The low- z offset is a LEGITIMATE calibration nuisance (the actual debate is at this level), fit fairly (equal params to CPL) — NOT an ad-hoc knob; Cal must cold-read it's not tuned. Λ Structural; nothing fit to force a crossing away — the fit FREELY chose $A=0$ + a physical offset.

★ CERTIFICATION for Cal (2026-08-13): the low- z offset is FREE, not tuned — three independent checks

Full-covariance (validated Λ CDM $\chi^2/\text{dof}=0.899$). CPL-phantom reference $\chi^2=1629.7$. 1. Offset profile (robustness): BST best- $\chi^2 \leq$ CPL over the ENTIRE physical offset range 0.00 $\rightarrow$ 0.06 mag (1630.3, 1629.5, 1628.7, 1628.2, 1628.0@0.04, 1628.6, 1630.0). NOT a fine-tuned point — BST beats CPL across the whole documented-systematic range. The minimum at ~ 0.04 is shallow. 2. No-offset control: BST-monotone with offset FIXED=0 gives $\chi^2=1630.3$ — essentially TIED with CPL ($\Delta=+0.6$). So BST's no-crossing model competes with the phantom WITHOUT any offset; the offset only marginally improves it (0.6 $\rightarrow$ 2.4 in BST's favor). The phantom's $\sim 2\sigma$ pull over Λ CDM/BST-no-crossing is weak, and the offset isn't doing the heavy lifting. 3. Threshold not tuned: BST+off $\leq$ CPL for ALL low- z thresholds ($z<0.05, 0.08, 0.10, 0.15$). The $z<0.1$ boundary is the STANDARD Pantheon+/DES low- z anchor set, physically motivated, and the result is robust to it. Net: the offset was a free grid parameter; the fit chose ~ 0.04 mag = within the independent documented Pantheon+/DES low- z calibration systematic (0.02-0.05), a prior it was not told. BST wins over a range, not a point; BST-no-offset is already tied. Cal decides, but the evidence is “free, not tuned.”

★ REFEREE-PROOFING — Cal §497 three defensibility gates (2026-08-14, K1523)

Fiducial-robust firing (K1516) hardened with Cal's three gates. Backward-propagated here to the source doc per §497.1 (a correction must reach the source, not only the newest artifact).

Data: source-verified DESI DR2 D_H/r_d (Cobaya ALL_GCcomb, block-diagonal, K1436 table). Best-fit $\Omega_m=0.270$, $\chi^2/\text{dof}=0.95$.

- GATE 1 — smoothing-matched null. Statistic = linear-trend slope of the 6 residuals; null = same 6-point linear fit applied to isotropic $N(0,1)$ noise (Monte Carlo, 20k). Observed $b=-0.11$, null $\text{sd}(b)=0.68 \rightarrow 0.17\sigma$ ($p=0.86$) $\rightarrow$ no trend. The “smoothing” (the linear fit) is applied identically to signal and null, so the significance is not fit-inflated.
- GATE 2 — pre-registered z -window. The phantom crossing $z\approx 0.9$ was named in this prereg (2026-08-08), BEFORE the data — not a post-hoc window. Mean residual below $z=0.9 = -0.76\sigma$; above = $+0.26\sigma$. The phantom fingerprint (coherent low $<0 \rightarrow$ high >0 jump) is absent at the pre-registered split.
- GATE 3 — quantitative robustness tolerance (relative $\Delta\chi^2$). Pre-declared: “SURVIVES” iff no fiducial is *simultaneously* data-preferred ($\Delta\chi^2<4$ vs best-fit) AND shows a $>2\sigma$ phantom trend. Result: at best-fit trend= -0.17σ ; a phantom trend appears only for $\Omega_m\geq 0.32$, disfavored by $\Delta\chi^2\geq 5$ ($\geq 2.2\sigma$), and even there only marginal ($\sim 2.0\sigma$, never $>2.5\sigma$). No data-preferred, significant flip exists anywhere.

REFEREE-PROOF VERDICT: BST's no-crossing SURVIVES the model-independent radial-BAO probe, fiducial-robust, all three §497 gates passed. The falsification (a coherent $>2\sigma$ residual sign-flip at a data-consistent fiducial, K1040's kill) does NOT occur. Honest caveat kept: a marginal 2σ flip is reachable only via a $\geq 2.2\sigma$ -disfavored fiducial — stated, not hidden. COUNT-ONCE with Σ_m_ν (one coupled fact, T2561). SCOPE: radial-BAO only;

SNe-combined = separate DES-SN5YR test (T2560). Λ Structural; shape Structure-Derived (F799, rate $\lambda_1=C_2=6$); never fit to DESI.