

K1329 — RE-RUN PASS = FAIL (again). The registry has a DEEP stale-“Proved” layer, not a 4-item punch-list — every scan finds more (Cal’s ~14 expanded rows + T1952/T2032/T2058/T2167). Go SYSTEMATIC: a registry HEADER disclaimer (the “Proved” column = derivation-depth, NOT referee-consensus — task #91) + reframe the top-level Millennium proof-CLAIM rows + fix FOUR internal contradictions (ruled below). Cal cold-reads.

Keeper (final gate + consistency, 2026-08-10 ~10:40 EDT). My re-run found the whack-a-mole is unwinnable row-by-row — the registry’s March–May layer is thick with stale “Proved,” and four rows contradict BST’s OWN current results. Incremental sweeps keep missing more. The honest fix is systematic. I rule the four contradictions here (mine) and recommend the systematic close. Cal cold-reads.

RE-RUN = FAIL — the pattern (why go systematic)

Grace closed the original 4 items cleanly; Cal expanded the reframe list (~14 rows) — and those are ALL still “Proved”/“Complete”/“Closed” (T91, T1146, T1170, T1465, T896, T1400, T1411, T1750–1752, T1798/1799, T972/993). My re-run found EVEN MORE the general grep caught (T1952 “over-determined system”, T2032 “quark sector ZERO free parameters PROVED”, T2058 cosmic-ladder “Proved”, T2167 “YM proof mass gap is m_p ”). Every scan finds a new layer. Row-by-row is unwinnable and keeps failing the PASS. The systematic fix: 1. ★ REGISTRY HEADER DISCLAIMER (task #91, the sustainable fix): add a prominent header — “*The ‘Proved’ column denotes internal derivation-DEPTH, NOT referee-consensus proof. Rigor is the D/PD/I/C/S tier. No entry claims a Clay-Millennium solution; the Millennium attempts are honestly tiered as Attempts (K940).*” This covers the hundreds of subsidiary “Proved” lemmas/identifications in ONE stroke, honestly. 2. Reframe the top-level Millennium proof-CLAIM rows (Cal’s expanded list — the ones a reviewer reads as “BST proves X”): ★ T91 “All Nine Millennium-Class Proofs are AC(0)” (the meta-claim), T1146 “YM Complete”, T1170 “YM Kill Chain Complete”, T1465 “BSD CLOSED”, T896/T1400/T1411 (YM), T1750–1752 (BSD), T1798/1799 (NS), T972/993 (YM R^4). → Attempt (K940). (*Grace executes Cal’s per-row calibration.*) 3. Fix the FOUR internal contradictions (ruled below). Then re-run. This is the honest scope — bigger than “4 items,” and worth doing right; a fast sweep would ship “YM Complete — Proved.”

★ THE FOUR INTERNAL CONTRADICTIONS (mine to reconcile — these are worse than over-claims; they contradict BST’s OWN results)

1. T569 “ $P \neq NP$ Linearization” contradicts the Curvature Principle (T1272/T1338). BST’s own result: $P \neq NP$ is the ONE Millennium problem that CANNOT be linearized — it IS the curvature obstruction (“you cannot linearize curvature”). So “ $P \neq NP$ Linearization” is a category error against the corpus. RULING: retire the linearization framing for $P \neq NP$; the row is corrected to “ $P \neq NP$ = the curvature obstruction (non-linearizable), Attempt” — consistent with T1272/T1338. (The K940 banner already added isn’t enough — the *framing* is the contradiction.)
2. T1399 “938 MeV Is Full Theory Mass Gap” + T2167 “YM proof mass gap is m_p (T187)” both conflate proton with the pure-YM gap — contradicts K1059/K1060. BST’s own walk-back: the pure-Yang–Mills mass gap is the glueball (~1.7 GeV), NOT the proton (938 MeV, a full-QCD-with-quarks bound state). RULING: retire “938 = pure-YM mass gap” in T1399 AND T2167 (and re-check T187); the pure-YM gap = glueball per K1059/K1060; the proton 938 is the full-QCD bound state, a different quantity.
3. T2032 “quark sector ZERO free parameters, PROVED, closes quark sector” contradicts the current flavor terminus (K1328/K1016). We JUST established the up-quark mass is the ONE genuinely soft number (no forced form, Identified, run backwards). Yet T2032 gives $m_u = c_3 \cdot N_c^2 / (\text{rank}^2 \cdot g) = 2.13 \text{ MeV}$ (1.4% off) as “PROVED zero-free.” RULING + open reconcile: at minimum, retag T2032 (1.4–3.5% off = Identified, NOT Proved; NOT “zero free parameters”). AND reconcile with Lyra: is $m_u = c_3 \cdot N_c^2 / (\text{rank}^2 \cdot g)$ a *forced form* (then K1328’s “soft” is too pessimistic — a real question worth reopening) or an *Identified fit* (then T2032’s “closed/zero-free” is the over-claim)? @Lyra/@Grace — reconcile the m_u status; the two can’t

both stand.

- (Also flag, lower priority: T1952's "over-determined system" language is the consistency-web trap again — temper to "self-consistent web"; not a Millennium blocker.)

Effect on the GO timeline (honest)

The flagship paper itself is clean (Cal hostile-read passed it) — this is the REGISTRY (a clone-facing knowledge-base file). But the repo IS the clone, so a corpus-miner hits it. So GO still waits on the systematic registry close (header + top-level reframe + the 4 contradictions). This is more than a day's incremental — but the header-disclaimer approach makes it tractable in one pass. Re-run after. (I owned "one punch-list from GO" as optimistic in K1328; it's now clearly a systematic-registry job — better found by the re-run than by a reviewer.)

Don't let the gate stall the week (Casey's steer: press on, investigate)

The registry close is Cal+Grace's lane and should NOT freeze the forward program. Run the roadmap in parallel (see the team prompt): the strong-sector / YM-mass-gap arc + the falsifier scrub. Casey's week: investigate → Friday/Saturday summary → Sunday (his 71st) his own investigate-items. Plenty of runway.

— Keeper, K1329, 2026-08-10. RE-RUN = FAIL: registry has a DEEP stale-"Proved" layer (Cal's ~14 + T1952/T2032/T2058/T2167); go SYSTEMATIC — (1) HEADER disclaimer ("Proved"=depth not referee-proof, task #91), (2) reframe top-level Millennium proof-claims (T91/T1146/T1170/T1465/T896/T1400/T1411/T1750-52/T1798-99/T972/993), (3) fix 4 contradictions: T569 ($P \neq NP$ non-linearizable per Curvature Principle T1272/T1338 — retire linearization framing), T1399+T2167 (pure-YM gap=glueball not proton-938 per K1059/K1060), T2032 (quark sector NOT zero-free — up-mass Identified/soft K1016/K1328; reconcile m_u form with Lyra). Then re-run. Flagship paper itself clean; GO waits on the registry. Roadmap runs in parallel (strong sector + falsifier scrub). Cal cold-reads. Nothing pushed.