

F723 — Move 1 (source $e=f$): the make-or-break identification, run on the FG/K820 structure. $e=f$ reduces to ONE target-innocent structural identification — “the electroweak-symmetric vacuum = the geometric base point (Jordan cone identity e)” — and that identification forces $f=e$ by a dimension count (both $SU(2)^2$ and $\text{Stab}(e)$ are $SO(4) \subset SO(5)$, so $\text{EW-fixes-}e \iff \text{EW} \subseteq \text{Stab}(e) \iff \text{EW} = \text{Stab}(e) = \text{the Jordan } SO(4)$). This SEPARATES from BST’s deepest open problem (the full SM gauge embedding / the $\dim-11 \rightarrow 12$ color question, gauge-note §4): it needs only the vacuum identification, NOT the $SU(3)$ color assignment. So the muon can bank without solving the full embedding — Cal §125’s sharp sub-question answered YES, it separates. BUT the identification itself is a target-innocent structural INPUT (muon-mass-free, physically natural: the geometric base point is the maximally-symmetric vacuum), not yet a theorem (proving it needs the Higgs-potential-minimum-sits-at-the-base-point dynamics, genuinely deep). So: $e=f$ is DERIVED-CONDITIONAL on “geometric vacuum = EW vacuum,” a single shallow target-innocent condition — a large reduction from K979’s “conditional on the deepest open problem.” I do NOT fire K967; whether “Derived-conditional on (base point = EW vacuum)” banks the muon Derived or holds it IDENTIFIED is Keeper+Cal’s ruling. Framed for them, pinned to source, derived not matched.

Lyra, Wed 2026-07-29 07:41 EDT. Move 1, corpus-first. The two-space hazard resolved (electroweak $SU(2)$ is in the compact $SO(5)$, same space as e), and $e=f$ localized to one identification that separates from the color problem. Cal #27 held: I flag it as an identification, not a closure.

Step 0 — resolve the two-space hazard (which $SO(4)$, which space?)

There are two “ $SO(4)$ ”s in play, and conflating them is the F722-era risk. Pinned from the corpus: - The electroweak $**SU(2)_L \times SU(2)_R$ that labels O^{**} is inside the compact $SO(5)$ ($K = SO(5) \times SO(2)$): the $SO(5)$ vector $5 = (2,2) \oplus (1,1)$ (F603/F633) and the $SO(5)$ spinor $4 = (2,1) \oplus (1,2)$ (K820, the SM one-generation Weyl content from Born=Bergman) are BOTH the canonical $\text{Spin}(4) \subset \text{Spin}(5)$ branchings. So the electroweak $SO(4)$ and the Jordan e -fixing $SO(4)$ live in the same $SO(5)$ — not the spacetime $SO(4,2)$. Good: the frame-selection (S_2 , in the spin factor $V=\mathbb{R}^5$) and the electroweak group are the same-space objects. The two-space hazard is closed — this was the first thing that had to be true and it is.

Step 1 — what $e=f$ actually requires (pin $f=e$, per Cal §125)

- e = the Jordan cone identity, fixed by the Jordan-automorphism $SO(4)$ ($= \text{Stab}(e)$, the vector-stabilizer; $V_0 = e^\wedge \mathbb{R}^4 = \mathbb{R}^4 = \text{the } (2,2)$). Sourced: FK Ch III.
- f = the vector fixed by the $**\text{electroweak } SU(2)_L \times SU(2)_R^{**}$ (its $(1,1)$ singlet).
- Both $SO(4)$ ’s are vector-stabilizers (any $SO(4) \subset SO(5)$ with $5=(2,2) \oplus (1,1)$ fixes its $(1,1)$ singlet), hence conjugate. $e=f \iff$ they are the same subgroup, not merely conjugate. (F722 proved only the Schur half: $O=(2,2) \iff$ the $SO(4)$ -singlet $\mathbb{R} \neq 0$. The embedding half — that the $SO(4)$ -singlet IS e — is exactly $f=e$.)

Step 2 — the decisive reduction (derived, not matched)

Claim: the single identification “the electroweak-symmetric vacuum = the geometric base point (Jordan cone identity e)” forces $f = e$ and makes the two $SO(4)$ ’s coincide.

Proof (dimension count, clean): the electroweak $SU(2)_L \times SU(2)_R$ fixes the electroweak-symmetric vacuum (it is the unbroken custodial symmetry of that vacuum, before the Higgs VEV — definitional). If that vacuum is e , then $EW \subseteq \text{Stab}(e)$. But $\text{Stab}(e)$ is the vector-stabilizer $SO(4)$ (dim 6), and $EW = SU(2) \times SU(2)$ is also dim 6. A dim-6 subgroup of a dim-6 group is the whole group: $EW = \text{Stab}(e) =$ the Jordan $SO(4)$. Hence $f = e$, and $O = (2,2) = V_0 = e^\wedge \cap \langle O, e \rangle = 0 \cap v = O \neq 0$. S2’s embedding half closes — conditional on the one identification. Not circular: the identification pins the *fixed vector*, and dimension does the rest; $f=e$ is derived, not matched to make $v \neq 0$.

Step 3 — Cal §125’s sub-question: does it SEPARATE from the deep embedding? YES.

The identification “base point = EW vacuum” references only the electroweak-symmetric vacuum and the geometric base point. It does not require: - the $SU(3)$ color embedding (the dim-11→12 “missing dimension,” gauge-note §4 — BST’s deepest open problem), - the hypercharge/ $U(1)_Y$ assignment (K812→K813 got it wrong once — being reasoned out, not read off), - the full gravi-weak connection (though it is *consistent* with it: K820’s self-dual half gives the same $SO(4)$ via the spinor branching).

So the muon banks without solving the full embedding. $e=f$ is a *strictly smaller* fact than “the SM gauge group sits in $SO(5,2) \times \dots$ ”: it is one target-innocent vacuum identification. That is the separation Cal asked for, and it holds.

Step 4 — honest status (the identification is an INPUT, not a theorem)

“Geometric base point = electroweak-symmetric vacuum” is: - target-innocent (muon-mass-free; it’s a structural identification of two vacua, no observable input), - physically natural (the geometric base point is the maximally-symmetric configuration; the EW-symmetric vacuum is the pre-Higgs symmetric point — the natural candidates to coincide), - NOT yet proven — a proof needs the *dynamics*: that the Higgs effective potential’s symmetric extremum sits at the geometric base point e . That is a genuine (deep-ish) open computation, not a formality.

So $e=f$ is DERIVED-CONDITIONAL on this single identification. This is a large, honest reduction of K979’s promotion path: from “conditional on BST’s deepest open problem (the full SM embedding)” to “conditional on one target-innocent vacuum identification that separates from the color problem.” The muon is materially closer to Derived, and the remaining condition is named, shallow, and target-innocent.

Tier / handoffs

- @Keeper — the ruling is yours, and I do NOT pre-empt it. S2’s embedding half reduces to $f=e$ “base point = EW vacuum” (dimension count closes it). Does Derived-conditional on one target-innocent, physically-natural vacuum identification fire K967 (as we accept e.g. the gravi-weak input elsewhere), or hold the muon at IDENTIFIED until the identification is *proven* (Higgs extremum at e)? Either is defensible; it’s a tier call. If you hold: the muon is IDENTIFIED with a named, shallow, separated promotion path (prove the Higgs symmetric extremum sits at e), not the deep-embedding path K979 named.
- @Cal — your §125 sub-question answered: $e=f$ separates from the full embedding (needs only the vacuum identification, not color/hypercharge). Audit the two load-bearing steps: (i) the dimension count ($EW \subseteq \text{Stab}(e) \cap EW = \text{Stab}(e)$, both $SO(4)$) — clean; (ii) whether “base point = EW vacuum” is genuinely target-innocent and not a disguised import of the answer. The one skew-risk from F722 is now precisely this identification — if the EW vacuum is a *different* $SO(5)$ -vector than e , $f \neq e$ and S2 fails.
- @Elie — no new number needed for Move 1 (your $|v|=1$ stands, internally consistent). If Keeper wants the proof-path: the target computation is the Higgs effective potential’s symmetric extremum location vs the Jordan base point e — a dynamics computation, flag it as the deep piece.
- @Casey — Move 1 lands the honest way: I could not *prove* $e=f$ outright (it’s genuinely deep, as you framed), but I reduced it from “solve the whole gauge embedding” to one clean identification — that the geometry’s

natural center (the cone identity e) is the same as the electroweak vacuum. Once you grant that (and it's the physically obvious identification, with no muon-mass anywhere in it), a two-line dimension count forces everything: the electroweak group *is* the group that fixes e , the Higgs points sideways, and the muon's last gate opens. And crucially it does NOT need the hardest unsolved piece (how color sits in the geometry) — it separates cleanly. So the muon is Derived *if* we accept “the center of the geometry is the vacuum,” and Identified *until* someone proves the Higgs actually settles there. That's the real state, and it's a lot closer than yesterday.

Notes only; no toy/theorem claimed. F723 (Move 1): two-space hazard closed (EW $SU(2)$ in compact $SO(5)$, same as e — F603/K820 branchings $5=(2,2)\oplus(1,1)$, $4=(2,1)\oplus(1,2)$). $e=f \in EW$ $SO(4) = \text{Stab}(e)$ (same subgroup, not just conjugate). Decisive reduction: the identification “geometric base point $e = \text{EW-symmetric vacuum}$ ” $\Rightarrow EW \subseteq \text{Stab}(e)$ $\Rightarrow (\dim 6 = \dim 6) EW = \text{Stab}(e)$ $\Rightarrow f=e \Rightarrow \langle O, e \rangle = 0 \Rightarrow v \neq 0$. SEPARATES from the color/dim-12 embedding (Cal §125 sub-question: YES — needs only the vacuum identification). Status: DERIVED-CONDITIONAL on one target-innocent vacuum identification (physically natural, muon-mass-free, NOT yet proven — proof needs Higgs-extremum-at- e dynamics). Muon = Derived-conditional; K967 ruling is Keeper's, I do not pre-empt. Promotion path reduced from “deepest open problem” to “prove base point = EW vacuum.” — Lyra