

How to Read a BST Claim: the Tier System in One Page

Keeper, v0.1, 2026-08-24 — packet document (R2). Scope: a physicist reading only this page can correctly predict what evidence stands behind any labeled claim in this packet. That sentence is this page’s own test: if you find a labeled claim whose support surprises you after reading this, the page has failed and we want to know.

The five tiers — what each label promises

DERIVED (D). The result follows from the geometry of D_{IV}^5 by an argument you can check, with no step chosen to reach the answer. Where a Derived claim consumes an input, the input is named *in the statement* — e.g., the hypercharge theorem carries its four inputs (one observational) inside its own sentence. Derived-conditional means exactly that: derived, given named conditions that are themselves stated, not hidden.

PARTIALLY DERIVED (PD) — always in explicit-split form. Part of the claim is forced and part is not, and the label tells you which is which: “*magnitude derived, branch data-picked*” or “*order derived, value input*.” A PD claim never blends the halves. If you see a PD label without its split, that is a defect — report it.

IDENTIFIED (I). A formula in BST’s integers matches a measured value, and we do not claim to know why. Identification is a candidate, not evidence: our own null tests show that expressibility in small integers is cheap (the fine-structure constant’s famous 137 sits in the *bulk* of what such forms can reach, and we say so). An Identified claim graduates only when a mechanism is derived — and several never will.

CANDIDATE (C). A mechanism or identity with a stated test it has not yet passed — or a result awaiting its second, independent verification. Candidates carry their kill-conditions with them.

STRUCTURAL (S). A statement about the framework’s shape (what exists, what is impossible) rather than a number. Example: “*the geometry cannot supply a length scale*” is itself a proved theorem — an honest openness, labeled as such.

Compound and local labels — the dialect map (one table, so no label surprises you)

The corpus is written by several hands over months; you will meet compound labels. Every one maps onto the five tiers, and this table is the packet’s official mapping:

label you may meet	reads as	meaning
Derived-full	D	all conditions discharged; nothing held
Derived-conditional	D (conditions named)	derived, given conditions stated in the theorem itself
CONDITIONAL-FORCED	D (conditions named)	same species: the relation is forced given a named, still-open condition
derived-synthesis	D, scope-limited	the identification is structural and audited, but is a <i>reading</i> of derived parts, not a new derivation
<i>flagship abstract’s</i> derived / supported / open / runner	D / C-or-I / (admission) / S-species	a coarse four-word scale used in one abstract; “open” is an admission, not a tier; “runner” = scale-dependent, no fixed number

If you meet a label not on this table or the five above, that is a defect — report it.

The three disciplines behind the labels

1. Pre-registration. Bars, win/loss conditions, and failure modes are frozen *before* computations run. Both of this program’s recent headline negatives — a chirality criterion failing, a mass-mechanism dying at its own test — were scored against bars written in advance, and are reported as failures by name. A register of falsifiers (packet document) lists what would kill each claim, including the ones currently passing.
2. Target-innocence. An input is admissible in a derivation only if it exists independently of the answer. Where an input is the observed answer itself (e.g., the observed charged fermion spectrum in the hypercharge theorem), it is declared *observational* — in the theorem statement.
3. One input, stated. Dimensionless structure is derived; the theory takes one measured number as its ruler and says so. Ratios are predictions; the absolute scale is calibration. This is the same posture as GR (which posits G) and QM (which posits $\hbar$), made explicit.

What the labels do NOT mean

- *Derived* does not mean “beyond question”: it means the argument is exhibited and its inputs are on its face. Attack the argument, not our confidence.
- A count of registry tags is not a verification statistic; only per-claim tiers are.
- Agreement with data quoted without a σ , or against an unnamed dataset, does not count as *evidence weighting* here. Scope of that standard, stated honestly: it is enforced in the falsifier register and in Forcing & Evidence’s Part-III case studies (pinned releases, e.g. NuFIT 6.0 / PDG 2024 Rev 12, σ against the stated band). The F&E summary tables still carry 193 of 197 rows without error bars — a known deficiency with a named owner (the data-layer landing pass, Grace), recorded in F&E itself — and the flagship quotes precisions, not σ ’s, for most of its 26 rows. Where you see a bare %, read it as a deviation report, not a weighted claim.

The one-line test, restated

Pick any labeled claim in this packet. From the label alone you should be able to say: *what would change my mind about this, and what did its authors already try*. If you cannot, the label — or this page — is wrong, and either is a bug we will fix in public.

— *The tier system is maintained by the program’s internal auditor. The SYSTEM itself and the packet’s headline assignments have survived adversarial review by a referee whose written role is to break them (his findings are cited where they bit). Per-row adversarial review of every table entry is NOT claimed — rows carry their own audit citations where they exist, and rows without one have exactly the support their tier label promises, no more. This page’s original version overclaimed that sentence; the referee caught it; this is the corrected form — which is the system working as described.*